

# Bagging-final-model

December 15, 2020

```
[1]: import sys
sys.path.append("../")
import pandas as pd
from ortho_lib3 import *
import os
import matplotlib.pyplot as plt
import numpy as np
import math
```

```
[2]: exercises = Exercises.load("../Pickle/
↳sliced_transformed_exercises_9_12_all_categories.pickle")
#exercises = exercises.drop_category(1)
exercises.df
```

```
[2]:      angle_left_shoulder_xz_max_AF  angle_left_shoulder_xz_max_RF  \
0                2.705811                1.200405
1                2.757520                1.381570
2                2.691818                1.372206
3                2.524043                0.687781
4                2.640875                1.189572
..                ...                ...
115               1.201943                0.428302
116               2.186121                0.580735
117               1.882206                0.808175
118               1.237203                0.822418
119               2.602110                1.114584

      angle_right_shoulder_xz_max_AF  angle_right_shoulder_xz_max_RF  \
0                2.880950                1.022868
1                2.693907                1.300333
2                2.597640                1.352112
3                2.581135                0.931032
4                2.606757                1.312657
..                ...                ...
115               2.347702                0.429440
116               2.082403                0.712122
117               2.064096                0.987518
```

118	2.171692	0.780059
119	2.614965	1.054871

	angle_left_shoulder_yz_max_AB	angle_right_shoulder_yz_max_AB	\
0	2.536825	2.785053	
1	2.729818	2.894459	
2	2.312296	2.432454	
3	2.606495	2.639872	
4	2.666926	2.706126	
..	...	...	
115	0.897281	2.104353	
116	2.649783	2.220724	
117	2.459425	2.348511	
118	1.519678	2.371818	
119	2.574107	2.385383	

	diff_x_wrist_std_EL	diff_x_wrist_std_AF	diff_x_wrist_std_RF	\
0	0.051140	0.045686	0.047964	
1	0.046985	0.032754	0.022710	
2	0.018752	0.043444	0.059608	
3	0.045662	0.032816	0.036892	
4	0.042752	0.048538	0.042671	
..	...	...	...	
115	0.131204	0.461456	0.041716	
116	0.022393	0.089273	0.157484	
117	0.055195	0.085915	0.043581	
118	0.113686	0.182995	0.066054	
119	0.037779	0.128872	0.047050	

	diff_x_elbow_std_EL	...	angular_acc_xz_elbow_r_mean_AF	\
0	0.066650	...	0.013454	
1	0.035020	...	0.009875	
2	0.035862	...	0.012321	
3	0.030882	...	0.008630	
4	0.015893	...	0.007462	
..	...	...	...	
115	0.038488	...	0.030021	
116	0.014435	...	0.036877	
117	0.025565	...	0.026531	
118	0.041324	...	0.055863	
119	0.029085	...	0.032568	

	angular_acc_xz_elbow_r_std_AF	angular_acc_xz_elbow_r_mean_RF	\
0	0.009997	0.009142	
1	0.007717	0.009636	
2	0.009420	0.010026	
3	0.007913	0.007300	

4	0.006183	0.007197
..	...	...
115	0.039972	0.009998
116	0.036002	0.044232
117	0.030848	0.030147
118	0.051294	0.031281
119	0.019488	0.023850

	angular_acc_xz_elbow_r_std_RF	angular_vel_yz_elbow_l_std_AB \
0	0.008593	0.035317
1	0.008744	0.017855
2	0.008610	0.028611
3	0.007061	0.021368
4	0.006843	0.026717
..	...	...
115	0.007187	0.013514
116	0.048294	0.066856
117	0.026745	0.042285
118	0.030357	0.047952
119	0.021694	0.057128

	angular_vel_yz_elbow_r_std_AB	angular_acc_yz_elbow_l_mean_AB \
0	0.034019	0.010569
1	0.019760	0.009327
2	0.027344	0.009416
3	0.021825	0.006564
4	0.023068	0.007782
..	...	...
115	0.040844	0.009843
116	0.062815	0.036638
117	0.048514	0.019478
118	0.058858	0.028998
119	0.066880	0.029559

	angular_acc_yz_elbow_l_std_AB	angular_acc_yz_elbow_r_mean_AB \
0	0.008540	0.009179
1	0.008234	0.008376
2	0.007159	0.008765
3	0.007230	0.007175
4	0.011197	0.007356
..	...	...
115	0.008446	0.017797
116	0.035517	0.042551
117	0.018303	0.019738
118	0.029493	0.032410
119	0.022851	0.030719

	angular_acc_yz_elbow_r_std_AB
0	0.008611
1	0.006633
2	0.007664
3	0.007337
4	0.005814
..	...
115	0.019753
116	0.033384
117	0.018876
118	0.037222
119	0.026552

[120 rows x 78 columns]

```
[3]: exp = Experiment(exercises, y_condition= lambda y: np.all([y != 'Category_1'],
    ↪axis=0))
print(exp.df.shape)
columns = exp.df.columns.to_numpy()
```

(120, 78)

```
[4]: # check scikit-learn version
import sklearn
print(sklearn.__version__)
```

0.23.2

```
[5]: from sklearn.datasets import make_regression
from sklearn.model_selection import cross_val_score, KFold, StratifiedKFold
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.ensemble import BaggingRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import classification_report
# define dataset
X = exp.df.values
y = exp.y
print(X)
```

```
[[2.70581125 1.2004053 2.88095043 ... 0.00854023 0.0091792 0.00861107]
 [2.75751986 1.38156992 2.69390747 ... 0.00823437 0.00837636 0.00663325]
 [2.69181792 1.37220563 2.59764029 ... 0.00715904 0.00876495 0.00766446]
 ...
 [1.88220638 0.80817506 2.06409626 ... 0.01830264 0.01973769 0.01887624]
 [1.23720302 0.82241828 2.17169177 ... 0.02949289 0.03240998 0.03722172]
 [2.60211036 1.11458353 2.61496538 ... 0.02285141 0.03071891 0.0265518 ]]
```

```

[6]: skf = StratifiedKFold(n_splits=5, random_state=None, shuffle=False)

df_scores = pd.DataFrame()

precision_list = []
recall_list = []
f1_score_list = []
accuracy_list = []
i=1
dtr = DecisionTreeRegressor()

for train_index, test_index in skf.split(X, y):
    # print("TRAIN:", train_index, "TEST:", test_index)
    Xtrain, Xtest = X[train_index], X[test_index]
    ytrain, ytest = y[train_index], y[test_index]
    regr = BaggingRegressor(dtr, n_estimators=100, max_samples=0.5)
    regr.fit(Xtrain, ytrain)
    ypred = regr.predict(Xtest)
    report = classification_report(ytest, ypred.round(), output_dict=True)
    recall = (report.get('weighted avg').get('recall'))
    precision = (report.get('weighted avg').get('precision'))
    f1_score = (report.get('weighted avg').get('f1-score'))
    accuracy = (report.get('accuracy'))

    precision_list.append(precision)
    recall_list.append(recall)
    f1_score_list.append(f1_score)
    accuracy_list.append(accuracy)

    feature_importances = np.mean([tree.feature_importances_ for tree in regr.
    estimators_], axis=0)
    #feature_scores = pd.Series(regr.feature_importances_, index=exp.df.
    columns).sort_values(ascending=False)
    #feature_scores = pd.Series(dtr.feature_importances_, index=exp.df.columns)
    feature_scores = pd.Series(feature_importances, index=exp.df.columns)
    df_scores['column' + str(i)] = feature_scores.values
    i=i+1

df_scores['average'] = df_scores.mean(axis=1)
df_scores['feature'] = exp.df.columns
print(f'precision: {np.mean(precision_list)} \nrecall: {np.mean(recall_list)} \n
f1_score: {np.mean(f1_score_list)} \naccuracy: {np.mean(accuracy_list)}')

```

```

precision: 0.9036934824434824
recall: 0.8916666666666666
f1_score: 0.8852599559916634

```

accuracy: 0.8916666666666666

```
[7]: df_scores[['feature', 'average']].sort_values(by='average', ascending=False).  
      ↪head(10).reset_index(drop=True)
```

```
[7]:
```

	feature	average
0	z_elbow_max_AF	0.139699
1	z_wrist_max_AB	0.102989
2	z_wrist_max_AF	0.093109
3	z_elbow_max_AB	0.068611
4	diff_y_wrist_std_AB	0.050212
5	x_wrist_max_EL	0.049581
6	diff_z_wrist_std_AB	0.029628
7	diff_x_shoulder_std_EL	0.027127
8	angle_left_shoulder_yz_max_AB	0.022151
9	acc_wrists_x_r_std_EL	0.018211

## 1 Reports


```
[8]: report = classification_report(ytest, ypred.round())  
      print(report)
```

	precision	recall	f1-score	support	
	0.0	1.00	0.80	0.89	5
	1.0	0.95	1.00	0.97	19
accuracy				0.96	24
macro avg	0.97	0.90	0.93		24
weighted avg	0.96	0.96	0.96		24


## 2 Confusion Matrix

```
[9]: from sklearn.metrics import confusion_matrix  
      cm = confusion_matrix(ytest, ypred.round())  
      cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual_↪  
      ↪Negative:0'],  
                               index=['Predict Positive:1', 'Predict Negative:↪  
      ↪0'])  
      sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='PiYG')
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7f04d2a90e80>
```



```
[10]: x_ax = range(len(ytest))
plt.scatter(x_ax, ytest, s=5, color="grey", label="original")
plt.plot(x_ax, ypred, lw=0.8, color="hotpink", label="predicted")
plt.yticks([0,1])
plt.legend()
plt.show()
```



### 3 Bomen printen

```
[11]: estimators = regr.estimators_  
len(estimators)
```

```
[11]: 100
```

```
[12]: #for bomen in estimators:  
#     print (bomen)
```

```
[13]: #from sklearn import tree  
#from matplotlib import pyplot as plt  
#for bomen in estimators[5:]:  
#    boom = bomen  
#    boom_score = boom.score(X,y)  
#    fig, axes = plt.subplots(nrows =1,ncols = 1,figsize = (5,5), dpi=300)  
#    tree.plot_tree(boom, filled=True, feature_names= exercises.df.columns)
```



## 4 Hyper parameter tuning

### 5 n\_estimators

```
[14]: skf = StratifiedKFold(n_splits=5, random_state=None, shuffle=False)

mean_precision_list = []
mean_recall_list = []
mean_accuracy_list=[]


n_estimators_list = [k for k in range(1,500,25)]

for n_estimators in n_estimators_list:
    precision_list = []
    recall_list = []
    accuracy_list=[]
    for train_index, test_index in skf.split(X, y):
        Xtrain, Xtest = X[train_index], X[test_index]
        ytrain, ytest = y[train_index], y[test_index]
        regr = BaggingRegressor(n_estimators=n_estimators, max_samples= 0.3)
        regr.fit(Xtrain, ytrain)
        ypred = regr.predict(Xtest)
        report = classification_report(ytest, ypred.round(), output_dict=True)
        recall = (report.get('weighted avg').get('recall'))
        precision = (report.get('weighted avg').get('precision'))
        accuracy = (report.get('accuracy'))

        precision_list.append(precision)
        recall_list.append(recall)
        accuracy_list.append(accuracy)

    mean_precision_list.append(np.mean(precision_list))
    mean_recall_list.append(np.mean(recall_list))
    mean_accuracy_list.append(np.mean(accuracy_list))

[15]: plt.plot(n_estimators_list, mean_precision_list, label='precision')
plt.plot(n_estimators_list, mean_recall_list, label='recall', marker = 'o')
plt.plot(n_estimators_list, mean_accuracy_list, label='accuracy')
plt.xlabel('n_estimators')
plt.ylabel('precision, recall, accuracy')
plt.legend()
plt.show()
```



## 6 max\_samples

```
[16]: skf = StratifiedKFold(n_splits=5, random_state=None, shuffle=False)

mean_precision_list = []
mean_recall_list = []
mean_accuracy_list = []

max_samples_list = [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0]

for max_samples in max_samples_list:
    precision_list = []
    recall_list = []
    accuracy_list = []
    for train_index, test_index in skf.split(X, y):
        Xtrain, Xtest = X[train_index], X[test_index]
        ytrain, ytest = y[train_index], y[test_index]
        regr = BaggingRegressor(n_estimators=175, max_samples = max_samples)
        regr.fit(Xtrain, ytrain)
        ypred = regr.predict(Xtest)
        report = classification_report(ytest, ypred.round(), output_dict=True)
        recall = (report.get('weighted avg').get('recall'))
        precision = (report.get('weighted avg').get('precision'))
```

```
accuracy = (report.get('accuracy'))

precision_list.append(precision)
recall_list.append(recall)
accuracy_list.append(accuracy)

mean_precision_list.append(np.mean(precision_list))
mean_recall_list.append(np.mean(recall_list))
mean_accuracy_list.append(np.mean(accuracy_list))
```

```
/opt/jupyterhub/anaconda/lib/python3.6/site-  
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
/opt/jupyterhub/anaconda/lib/python3.6/site-  
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
/opt/jupyterhub/anaconda/lib/python3.6/site-  
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
/opt/jupyterhub/anaconda/lib/python3.6/site-  
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

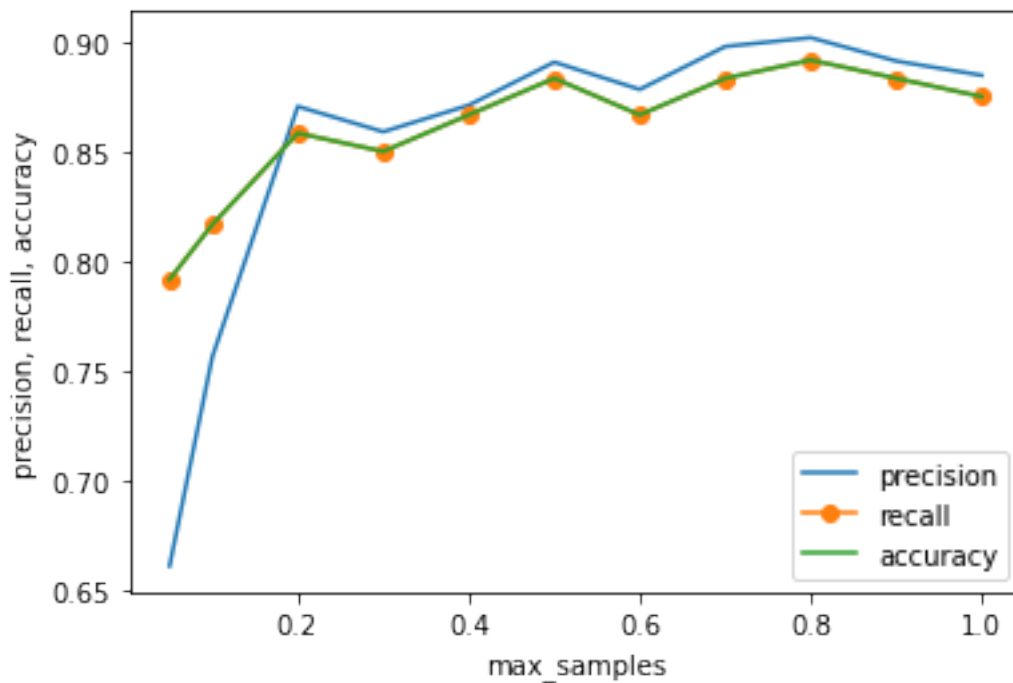
```
/opt/jupyterhub/anaconda/lib/python3.6/site-  
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
/opt/jupyterhub/anaconda/lib/python3.6/site-  
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
[17]: plt.plot(max_samples_list, mean_precision_list, label='precision')
plt.plot(max_samples_list, mean_recall_list, label='recall', marker= 'o')
plt.plot(max_samples_list, mean_accuracy_list, label='accuracy')
plt.xlabel('max_samples')
plt.ylabel('precision, recall, accuracy')
plt.legend()
plt.show()
```



```
[ ]:
```